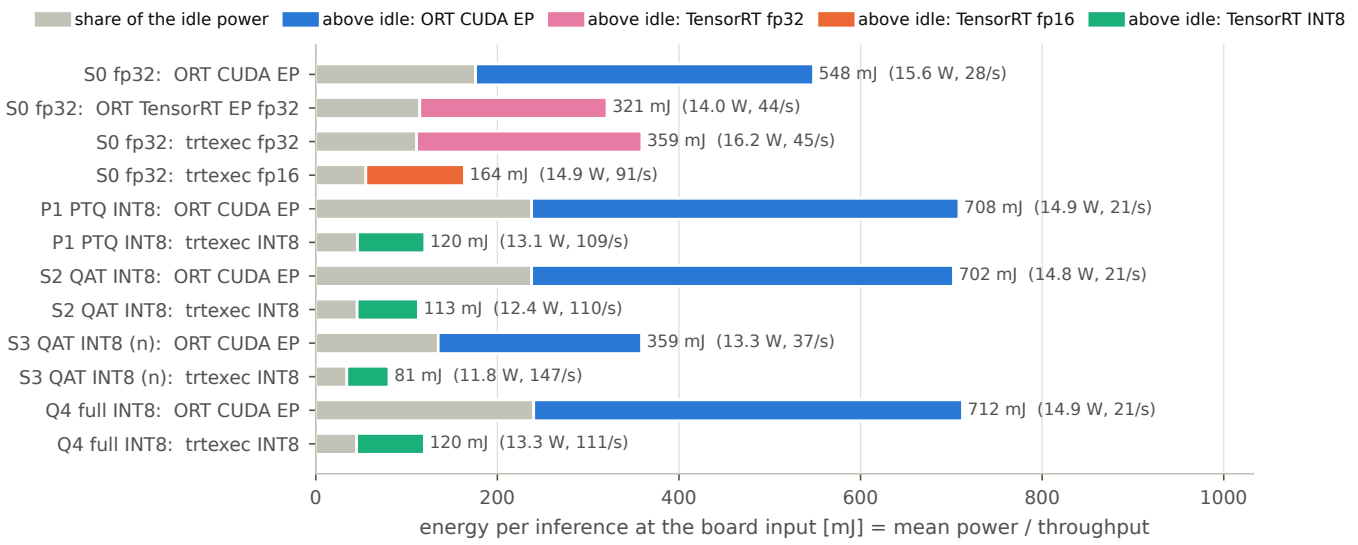
share of the idle power
above idle: ORT CUDA EP
above idle: TensorRT fp32
above idle: TensorRT fp16
above idle: TensorRT INT8
S0 fp32: ORT CUDA EP
S0 fp32: ORT TensorRT EP fp32
S0 fp32: trtexec fp32
S0 fp32: trtexec fp16
P1 PTQ INT8: ORT CUDA EP
P1 PTQ INT8: trtexec INT8
S2 QAT INT8: ORT CUDA EP
S2 QAT INT8: trtexec INT8
S3 QAT INT8 (n): ORT CUDA EP
S3 QAT INT8 (n): trtexec INT8
Q4 full INT8: ORT CUDA EP
Q4 full INT8: trtexec INT8
548 mJ (15.6 W, 28/s)
321 mJ (14.0 W, 44/s)
359 mJ (16.2 W, 45/s)
164 mJ (14.9 W, 91/s)
708 mJ (14.9 W, 21/s)
120 mJ (13.1 W, 109/s)
702 mJ (14.8 W, 21/s)
113 mJ (12.4 W, 110/s)
359 mJ (13.3 W, 37/s)
81 mJ (11.8 W, 147/s)
712 mJ (14.9 W, 21/s)
120 mJ (13.3 W, 111/s)
0
200
400
600
800
1000
energy per inference at the board input [mJ] = mean power / throughput